

TSSL Lab 3 - Nonlinear state space models and Sequential Monte Carlo

In this lab we will make use of a non-linear state space model for analyzing the dynamics of SARS-CoV-2, the virus causing covid-19. We will use an epidemiological model referred to as a Susceptible-Exposed-Infectious-Recovered (SEIR) model. It is a stochastic adaptation of the model used by the The Public Health Agency of Sweden for predicting the spread of covid-19 in the Stockholm region early in the pandemic, see [Estimates of the peak-day and the number of infected individuals during the covid-19 outbreak in the Stockholm region, Sweden February – April 2020](#).

The background and details of the SEIR model that we will use are available in the document *TSSL Lab 3 Predicting Covid-19 Description of the SEIR model* on LISAM. Please read through the model description before starting on the lab assignments to get a feeling for what type of model that we will work with.

DISCLAIMER

Even though we will use a type of model that is common in epidemiological studies and analyze real covid-19 data, you should *NOT* read too much into the results of the lab. The model is intentionally simplified to fit the scope of the lab, it is not validated, and it involves several model parameters that are set somewhat arbitrarily. The lab is intended to be an illustration of how we can work with nonlinear state space models and Sequential Monte Carlo methods to solve a problem of practical interest, but the actual predictions made by the final model should be taken with a big grain of salt.


We load a few packages that are useful for solving this lab assignment.

```
In [1]: import pandas # Loading data / handling data frames
import numpy as np
import matplotlib.pyplot as plt
plt.rcParams["figure.figsize"] = (10,6) # Increase default size of plots
```

3.1 A first glance at the data

The data that we will use in this lab is a time series consisting of daily covid-19-related intensive care cases in Stockholm from March 2020 to March 2021. As always, we start by loading and plotting the data.

```
In [2]: data=pandas.read_csv('SIR_Stockholm.csv',header=0)
y_sthlm = data['ICU'].values
u_sthlm = data['Date'].values
ndata = len(y_sthlm)
plt.plot(u_sthlm,y_sthlm)
plt.xticks(range(0, ndata, 14), u_sthlm[::14], rotation = 90) # Show only one tick per 2 week for clarity
plt.xlabel('Date')
plt.ylabel('New ICU cases')
plt.show()
```



Q0: What type of values can the observations y_t take? Is a Gaussian likelihood model a good choice if we want to respect the properties of the data?

A0: y_t can take non-negative integer values. A Gaussian likelihood model is not a good choice if we want to respect the properties of the data because the data is discrete not continuous.

3.2 Setting up and simulating the SEIR model

In this section we will set up a SEIR model and use this to simulate a synthetic data set. You should keep these simulated trajectories, we will use them in the following sections.

```
In [3]: from tssltools_lab3 import Param, SEIR

        """For Stockholm the population is probably roughly 2.5 million."""
        population_size = 2500000

        """ Binomial probabilities (p_se, p_ei, p_ir, and p_ic) and the transmission rate (rho)"""
        pse = 0          # This controls the rate of spontaneous s->e transitions. It is set to zero for this lab.
        pei = 1 / 5.1    # Based on FHM report
        pir = 1 / 5      # Based on FHM report
        pic = 1 / 1000   # Quite arbitrary!
        rho = 0.3        # Quite arbitrary!

        """ The instantaneous contact rate b[t] is modeled as
            b[t] = exp(z[t])
            z[t] = z[t-1] + epsilon[t], epsilon[t] ~ N(0,sigma_epsilon^2)
        """
        sigma_epsilon = .1

        """ For setting the initial state of the simulation"""
        i0 = 1000 # Mean number of infectious individuals at initial time point
        e0 = 5000 # Mean number of exposed...
        r0 = 0    # Mean number of recovered
        s0 = population_size - i0 - e0 - r0 # Mean number of susceptible
        init_mean = np.array([s0, e0, i0, 0.], dtype=np.float64) # The last 0. is the mean of z[0]

        """All the above parameters are stored in params."""
        params = Param(pse, pei, pir, pic, rho, sigma_epsilon, init_mean, population_size)
```

```
""" Create a model instance"""
model = SEIR(params)
```

Q1: Generate 10 different trajectories of length 200 from the model and plot them in one figure. Do the trajectories look reasonable? Could the data have been generated using this model?

For reproducibility, we set the seed of the random number generator to 0 before simulating the trajectories using `np.random.seed(0)`

Save these 10 generated trajectories for future use.

(hint: The SEIR class has a simulate method)

```
In [4]: np.random.seed(0)
        help(model.simulate)
```

Help on method simulate in module tssltools_lab3:

simulate(T, N=1) method of tssltools_lab3.SEIR instance

Simulates the SEIR model for a given number of time steps. Multiple trajectories can be simulated simultaneously.

Parameters

T : int

Number of time steps to simulate the model for.

N : int, optional

Number of independent trajectories to simulate. The default is 1.

Returns

alpha : ndarray

Array of size (d,N,T) with state trajectories. `alpha[:,i,:]` is the *i*:th trajectory.

y : ndarray

Array of size (1,N,T) with observations.


```
In [5]: np.random.seed(0)
        trajectory_length = 200
        trajectory_number = 10
        alpha, y = model.simulate(T=trajectory_length, N=trajectory_number)
```

```
print(y.shape)
for i in range(trajectory_number):
    plt.plot(y[0,i,:], label = f'Trajectory{i+1}')

plt.xlabel("Length")
plt.ylabel("Observations: New ICU Cases")
plt.title("10 Different Trajectories Of Length 200 ")
plt.legend()
```

(1, 10, 200)

Out[5]: <matplotlib.legend.Legend at 0x16be8e89b90>



A1: The trajectories seem unreasonable when we compared the real data. The value of ICU Cases in real data can go up to around 25-30, however in trajectories it is 500-600. So there is a significant discrepancy between real data and trajectories. Therefore the data could not have been generated using this model.

3.3 Sequential Importance Sampling

Next, we pick out one trajectory that we will use for filtering. We use simulated data to start with, since we then know the true underlying SEIR states and can compare the filter results with the ground truth.

Q2: Implement the **Sequential Importance Sampling** algorithm by filling in the following functions.

The **exp_norm** function should return the normalized weights and the log average of the unnormalized weights. For numerical reasons, when calculating the weights we should "normalize" the log-weights first by removing the maximal value.

Let $\bar{\omega}_t = \max(\log \omega_t^i)$ and take the exponential of $\log \tilde{\omega}_t^i = \log \omega_t^i - \bar{\omega}_t$. Normalizing $\tilde{\omega}_t^i$ will yield the normalized weights!

For the log average of the unnormalized weights, care has to be taken to get the correct output, $\log(1/N \sum_{i=1}^N \tilde{\omega}_t^i) = \log(1/N \sum_{i=1}^N \omega_t^i) - \bar{\omega}_t$. We are going to need this in the future, so best to implement it right away.

(hint: look at the SEIR model class, it contains all necessary functions for propagation and weighting)

```
In [6]: from tssltools_lab3 import smc_res

def exp_norm(logwgt):
    """
    Exponentiates and normalizes the log-weights.

    Parameters
    -----
    logwgt : ndarray
        Array of size (N,) with log-weights.

    Returns
    -----
    wgt : ndarray
        Array of size (N,) with normalized weights, wgt[i] = exp(logwgt[i])/sum(exp(logwgt)),
        but computed in a /numerically robust way/!
    logZ : float
        log of the normalizing constant, logZ = log(sum(exp(logwgt))),
        but computed in a /numerically robust way/!
    """

    logwgt_max = np.max(logwgt) # maximum Log-weight
    wgt_unnormalized = np.exp(logwgt - logwgt_max) # unnormalized weights
```



```

wgt_sum = np.sum(wgt_unnormalized)
wgt = wgt_unnormalized / wgt_sum # normalize weights
logZ = logwgt_max + np.log(wgt_sum) # log of normalizing constant

return wgt, logZ

def ESS(wgt):
    """
    Computes the effective sample size.

    Parameters
    -----
    wgt : ndarray
        Array of size (N,) with normalized importance weights.

    Returns
    -----
    ess : float
        Effective sample size.
    """
    ess = (np.sum(wgt**2)/np.sum(wgt))
    return ess

def sis_filter(model, y, N):
    d = model.d
    n = len(y)

    # Allocate memory
    particles = np.zeros((d, N, n), dtype = float) # All generated particles
    logW = np.zeros((1, N, n)) # Unnormalized log-weight
    W = np.zeros((1, N, n)) # Normalized weight
    alpha_filt = np.zeros((d, 1, n)) # Store filter mean
    N_eff = np.zeros(n) # Efficient number of particles
    logZ = 0. # Log-likelihood estimate

    # Filter Loop
    for t in range(n):
        # Sample from "bootstrap proposal"
        if t == 0:
            particles[:, :, 0] = model.sample_state(alpha0 = None, N = N) # Initialize from p(alpha_1)
            logW[0, :, 0] = model.log_lik(y=y[0], alpha=particles[:, :, 0]) # Compute weights

```

```

else:
    particles[:, :, t] = model.sample_state(alpha0 = particles[:, :, t-1], N = N )# Propagate according to dynamics
    logW[0, :, t] = logW[0, :, t-1]+model.log_lik(y=y[t], alpha=particles[:, :, t]) # Update weights

    # Normalize the importance weights and compute N_eff
    W[0, :, t], _ = exp_norm(logW[0, :, t])
    N_eff[t] = ESS(W[0, :, t])

    # Compute filter estimates
    alpha_filt[:, 0, t] = np.sum(particles[:, :, t] * W[0, :, t], axis = 1)

return smc_res(alpha_filt, particles, W, logW=logW, N_eff=N_eff)

```

Q3: Choose one of the simulated trajectories and run the SIS algorithm using $N = 100$ particles. Show plots comparing the filter means from the SIS algorithm with the underlying truth of the Infected, Exposed and Recovered.

Also show a plot of how the ESS behaves over the run.

(hint: In the model we use the S, E, I as states, but S will be much larger than the others. To calculate R , note that $S + E + I + R = \text{Population}$)

```

In [7]: N = 100

trajectory = 1

sis = sis_filter(model = model, y = y[0, trajectory, :], N = N)

R_true = population_size - np.sum(alpha[:, 3, trajectory, :], axis = 0)

R_sis = population_size - np.sum(sis.alpha_filt[:, 3, 0, :], axis = 0)

fig, ax = plt.subplots(2, 2, figsize=(12, 8))

# Infected
ax[0, 0].plot(alpha[2, trajectory, :], label='Underlying Truth', color='red')
ax[0, 0].plot(sis.alpha_filt[2, 0, :], label='Filter Mean', color='purple')
ax[0, 0].set(title='Infected', xlabel='Time', ylabel='Number of Infected')
ax[0, 0].legend()
ax[0, 0].grid(True)

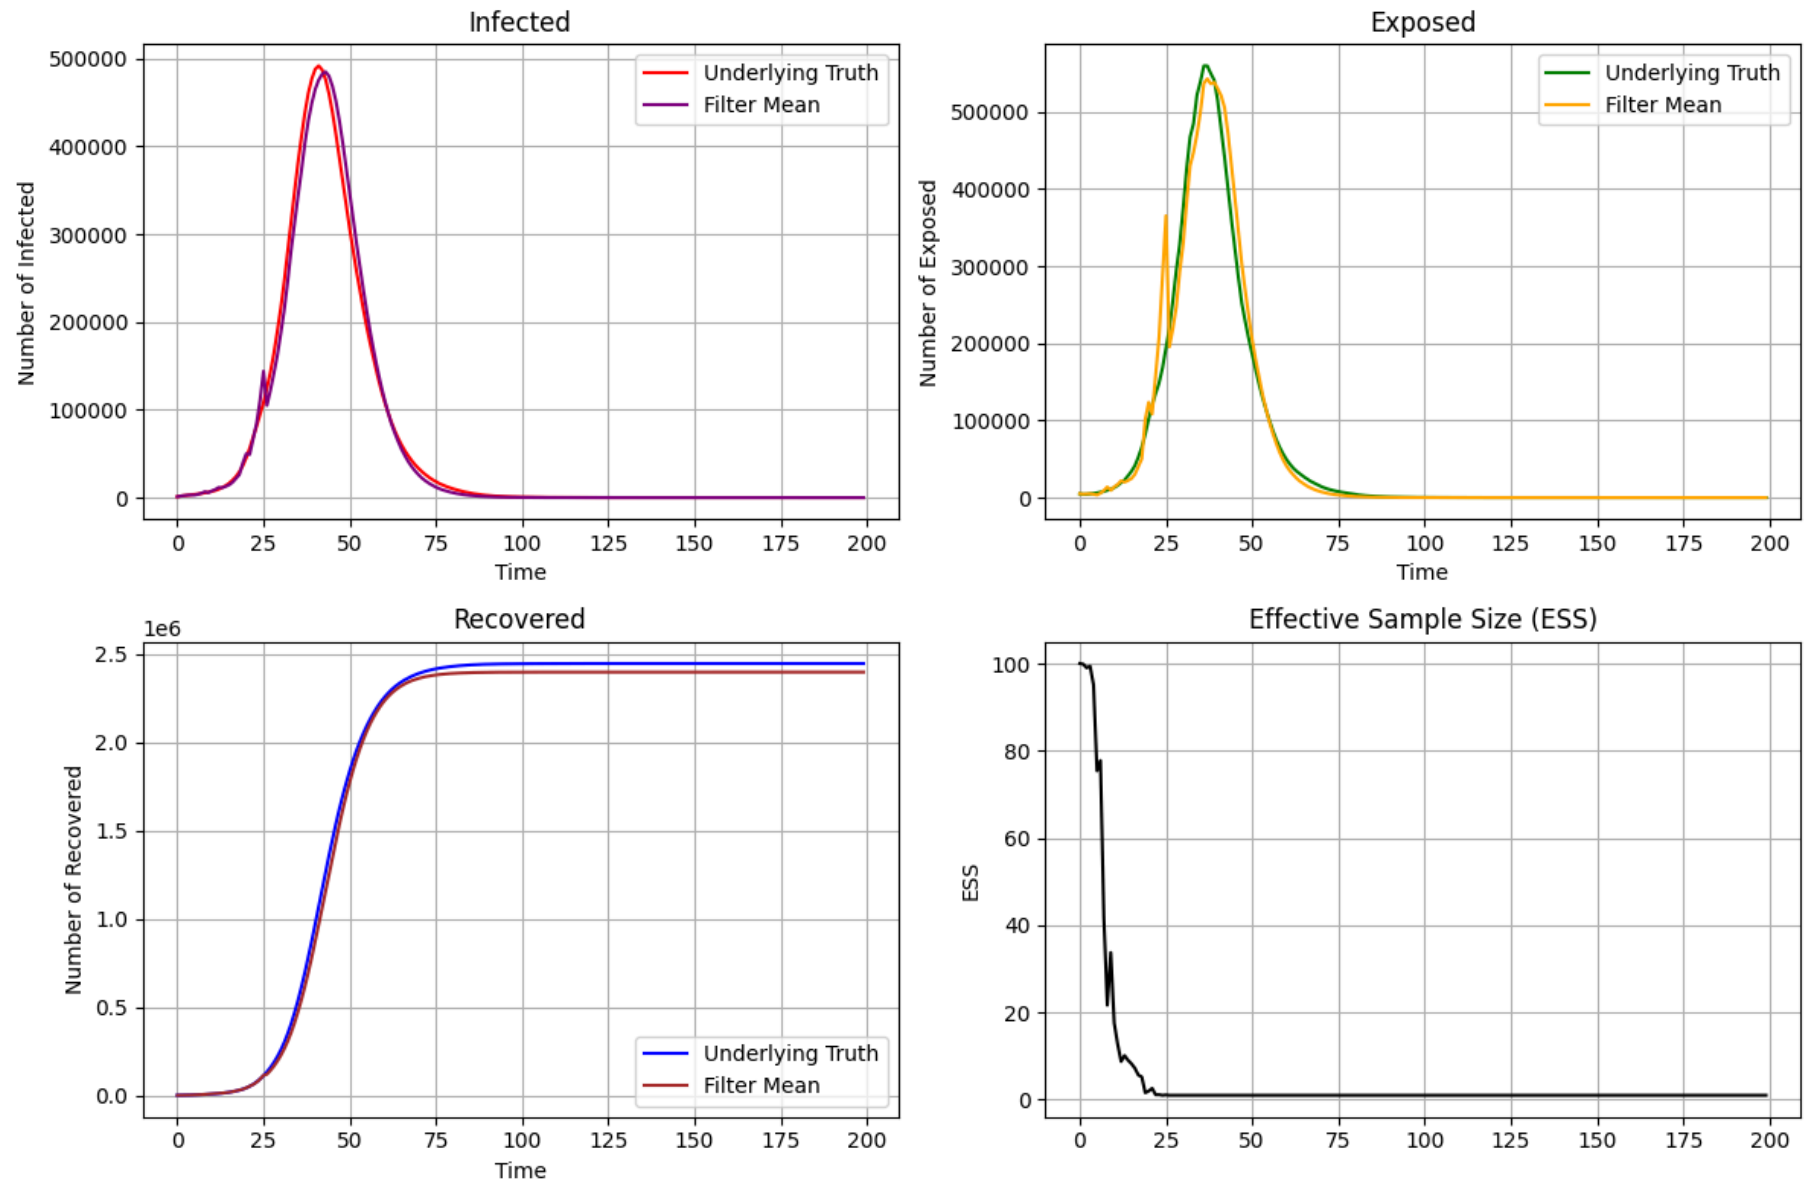
```

```
# Exposed
ax[0, 1].plot(alpha[1, trajectory, :], label='Underlying Truth', color='green')
ax[0, 1].plot(sis.alpha_filt[1, 0, :], label='Filter Mean', color='orange')
ax[0, 1].set(title='Exposed', xlabel='Time', ylabel='Number of Exposed')
ax[0, 1].legend()
ax[0, 1].grid(True)

# Recovered
ax[1, 0].plot(R_true, label='Underlying Truth', color='blue')
ax[1, 0].plot(R_sis, label='Filter Mean', color='brown')
ax[1, 0].set(title='Recovered', xlabel='Time', ylabel='Number of Recovered')
ax[1, 0].legend()
ax[1, 0].grid(True)

# Effective Sample Size
ax[1, 1].plot(sis.N_eff, color='black')
ax[1, 1].set(title='Effective Sample Size (ESS)', ylabel='ESS')
ax[1, 1].grid(True)

plt.tight_layout()
plt.show()
```



3.4 Sequential Importance Sampling with Resampling

Pick the same simulated trajectory as for the previous section.

Q4: Implement the **Sequential Importance Sampling with Resampling** or **Bootstrap Particle Filter** by completing the code below.

```
In [8]: def bpf(model, y, numParticles):
    d = model.d
    n = len(y)
    N = numParticles

    # Allocate memory
    particles = np.zeros((d, N, n), dtype = float) # All generated particles
    logW = np.zeros((1, N, n)) # Unnormalized log-weight
    W = np.zeros((1, N, n)) # Normalized weight
    alpha_filt = np.zeros((d, 1, n)) # Store filter mean
    N_eff = np.zeros(n) # Efficient number of particles
    logZ = 0. # Log-likelihood estimate

    # Filter loop
    for t in range(n):
        # Sample from "bootstrap proposal"
        if t == 0: # Initialize from prior
            particles[:, :, 0] = model.sample_state(N = N)
        else: # Resample and propagate according to dynamics
            ind = np.random.choice(N, N, replace=True, p=W[0, :, t-1])
            resampled_particles = particles[:, ind, t-1]
            particles[:, :, t] = model.sample_state(alpha0 = resampled_particles, N=N)

        # Compute weights
        logW[0, :, t] = model.log_lik(y[t], particles[:, :, t])
        W[0, :, t], logZ_now = exp_norm(logW[0, :, t])
        logZ += logZ_now - np.log(N) # Update log-likelihood estimate
        N_eff[t] = ESS(W[0, :, t])

        # Compute filter estimates
        alpha_filt[:, 0, t] = np.sum(W[0, :, t] * particles[:, :, t], axis=1)

    return smc_res(alpha_filt, particles, W, N_eff = N_eff, logZ = logZ)
```

Q5: Use the same simulated trajectory as above and run the BPF algorithm using $N = 100$ particles. Show plots comparing the filter means from the Bootstrap Particle Filter algorithm with the underlying truth of the Infected, Exposed and Recovered. Also show a plot of how the ESS

behaves over the run. Compare this with the results from the SIS algorithm.

```
In [9]: N = 100

trajectory = 1

bpf2 = bpf(model = model, y = y[0, trajectory, :], numParticles = N)

R_true_bpf = population_size - np.sum(alpha[:,3, trajectory, :], axis = 0)

R_bpf = population_size - np.sum(bpf2.alpha_filt[:,3, 0, :], axis = 0)

fig, ax = plt.subplots(2, 2, figsize=(12, 8))

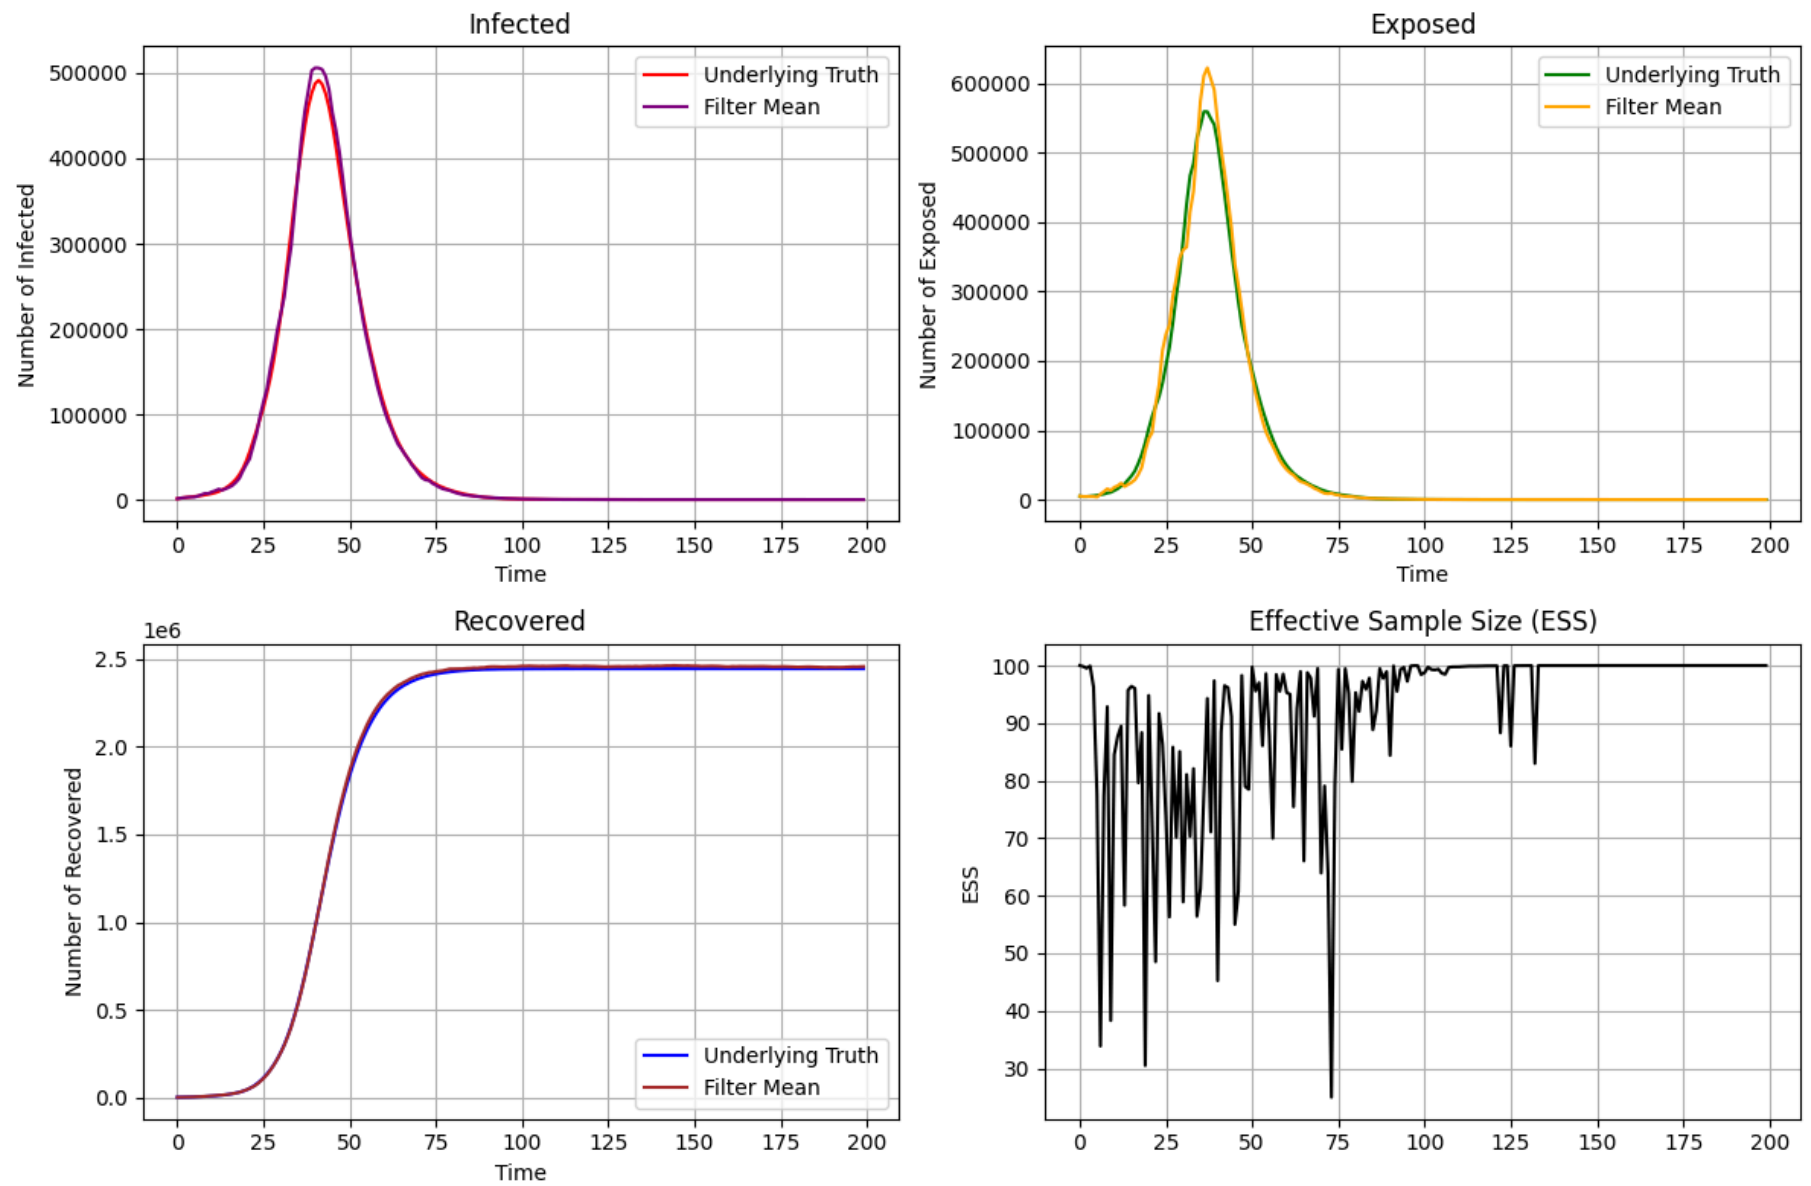
# Infected
ax[0, 0].plot(alpha[2, trajectory, :], label='Underlying Truth', color='red')
ax[0, 0].plot(bpf2.alpha_filt[2, 0, :], label='Filter Mean', color='purple')
ax[0, 0].set(title='Infected', xlabel='Time', ylabel='Number of Infected')
ax[0, 0].legend()
ax[0, 0].grid(True)

# Exposed
ax[0, 1].plot(alpha[1, trajectory, :], label='Underlying Truth', color='green')
ax[0, 1].plot(bpf2.alpha_filt[1, 0, :], label='Filter Mean', color='orange')
ax[0, 1].set(title='Exposed', xlabel='Time', ylabel='Number of Exposed')
ax[0, 1].legend()
ax[0, 1].grid(True)

# Recovered
ax[1, 0].plot(R_true_bpf, label='Underlying Truth', color='blue')
ax[1, 0].plot(R_bpf, label='Filter Mean', color='brown')
ax[1, 0].set(title='Recovered', xlabel='Time', ylabel='Number of Recovered')
ax[1, 0].legend()
ax[1, 0].grid(True)

# Effective Sample Size
ax[1, 1].plot(bpf2.N_eff, color='black')
ax[1, 1].set(title='Effective Sample Size (ESS)', ylabel='ESS')
ax[1, 1].grid(True)
```

```
plt.tight_layout()  
plt.show()
```



A5: When we compare the output of Sequential Importance Sampling (SIS) algorithm and Bootstrap Particle Filter (BPF) in Infected, the filter mean in SIS in the beginning there are some sharp and sudden line which in BPF there is not. After a while there is only a little deviation around peak but it is very little and we can say that filter mean can track the underlying truth. However, in BPA this slight deviation is less than SIS, so BPF could be slightly more accurate. For the Exposed, we can see that in SIS the filter mean have some struggle to track the underlying truth in the begining as can be seen sudden increses and decreases and also at the peak filter mean is a little under the underlying truth. However for Exposed in BPF, filter mean does not have sharp decreasing and increasing at the begining and also track the underlying truth well. It can be seen that the filter mean is slightly over the truth at the peak and it causes a less deviation. But the BPF filter mean is visibly more aligned with the underlying truth. For Recovered in both SIS and BPF filter mean track the truth very closely and in BPF it seems filter mean like a bit overlapping the truth. For Effective Simple Size graph, in SIS ESS decreases rapidly and fastly but after a while it goes around near zero. The reason of falling could be because of the weight degeneracy in the SIS algorithm. In BPF, the ESS start to falling and incresing until a certain time but after that time it stabilizes to 100. So BPF mitigates the falling rapidly and quickly issue like in SIS by resampling paricles which is not used in SIS algorithm. With resampling particles in BPF, we have more balanced weight distribution.

3.5 Estimating the data likelihood and learning a model parameter

In this section we consider the real data and learning the model using this data. For simplicity we will only look at the problem of estimating the ρ parameter and assume that others are fixed.

You are more than welcome to also study the other parameters.

Before we begin to tweak the parameters we run the particle filter using the current parameter values to get a benchmark on the log-likelihood.

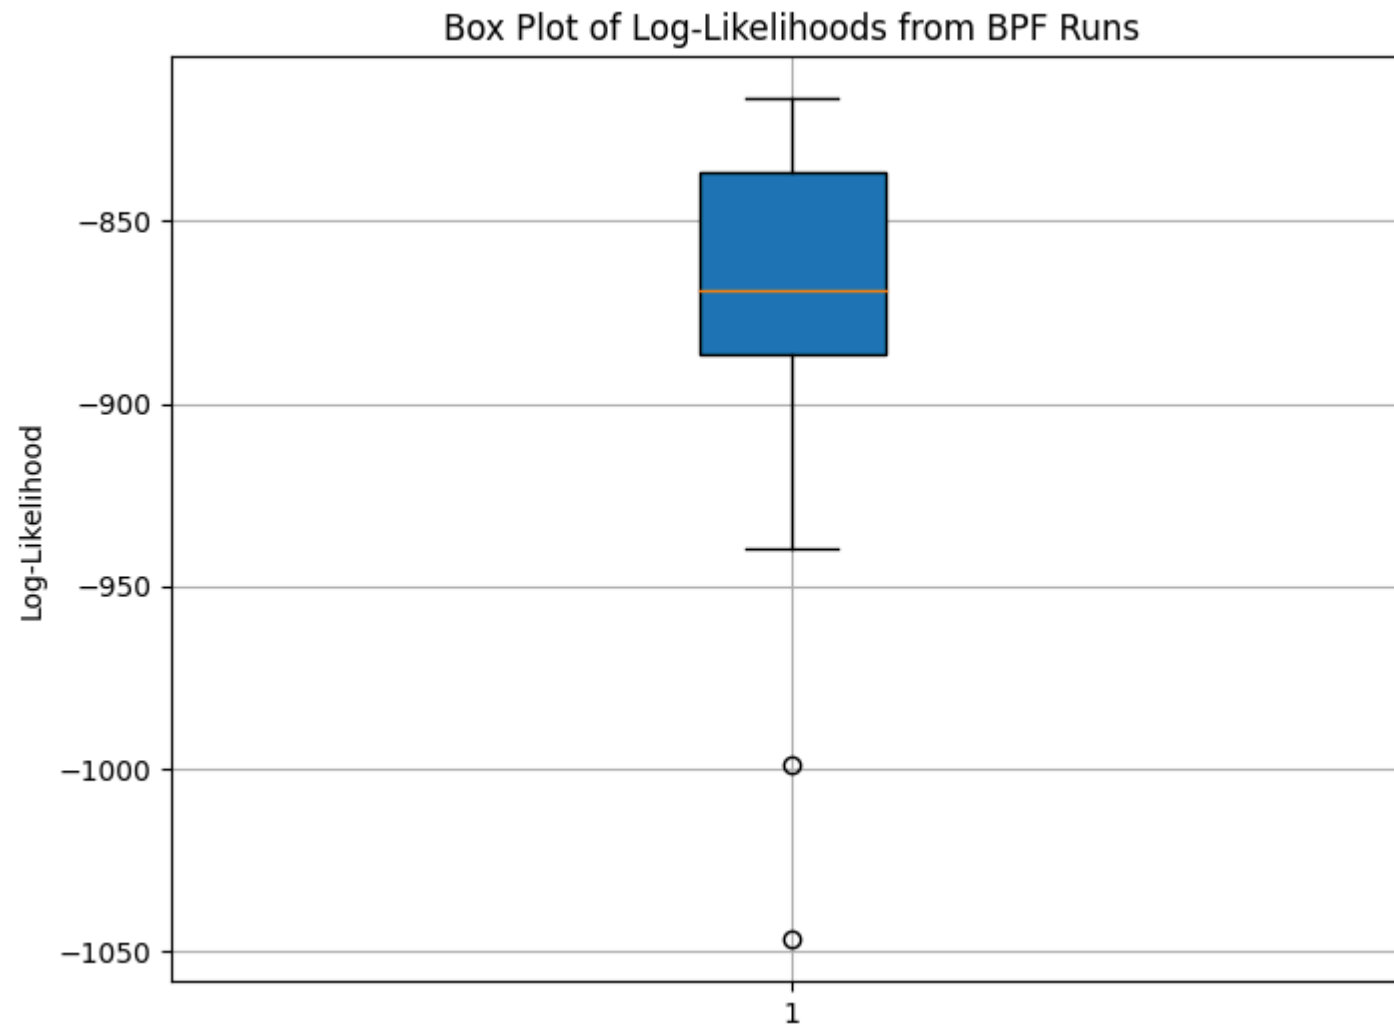
Q6: Run the bootstrap particle filter using $N = 200$ particles on the real dataset and calculate the log-likelihood. Rerun the algorithm 20 times and show a box-plot of the log-likelihood.

```
In [10]: N = 200
times = 20
log_likelihoods = []

for i in range(times):
    bpf_results = bpf(model=model, y=y_sthlm, numParticles=N)
    log_likelihoods.append(bpf_results.logZ)
```



```
plt.figure(figsize=(8, 6))
plt.boxplot(log_likelihoods, patch_artist=True)
plt.title('Box Plot of Log-Likelihoods from BPF Runs')
plt.ylabel('Log-Likelihood')
plt.grid(True)
plt.show()
```



Q7: Make a grid of the ρ parameter in the interval $[0.1, 0.9]$. Use the bootstrap particle filter to calculate the log-likelihood for each value. Run the bootstrap particle filter using $N = 200$ multiple times (at least 20) per value and use the average as your estimate of the log-likelihood. Plot the log-likelihood function and mark the maximal value.

(hint: use `np.linspace` to create a grid of parameter values)

```
In [13]: rho = np.linspace(0.1, 0.9, 20)
N = 200
times = 42
log_likelihoods2 = np.zeros((len(rho), times))

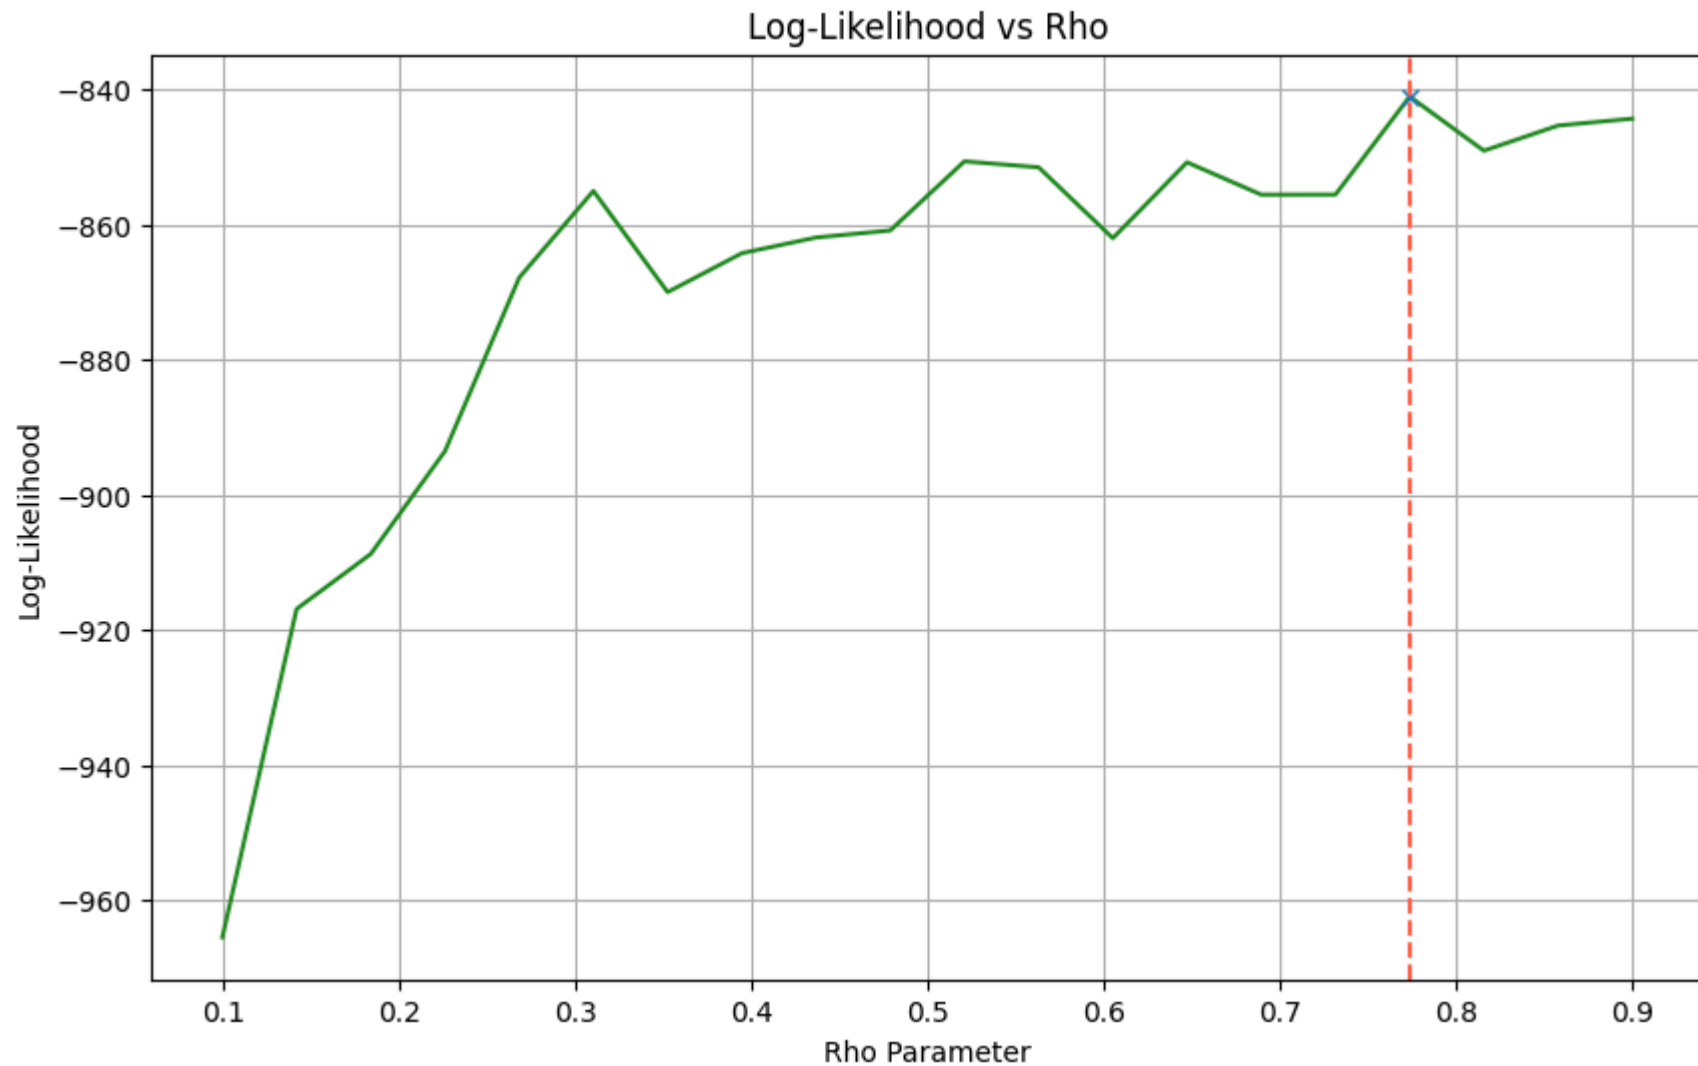
for i in range(len(rho)):
    model.set_param(rho = rho[i])

    for j in range(times):
        bpf_results2 = bpf(model = model, y = y_sthlm, numParticles = N)
        log_likelihoods2[i,j] = bpf_results2.logZ

mean = np.mean(log_likelihoods2, axis = 1)

plt.figure(figsize=(10, 6))
plt.plot(rho, mean, color = "forestgreen")
plt.plot(rho[mean == max(mean)], max(mean), marker = 'x')
plt.axvline(x = rho[mean == max(mean)], color = "tomato", ls = '--')
plt.title('Log-Likelihood vs Rho')
plt.xlabel('Rho Parameter')
plt.ylabel('Log-Likelihood')
plt.grid(True)
plt.show()

print("Optimal rho is:", rho[mean == max(mean)])
```



Optimal rho is: [0.77368421]

Q8: Run the bootstrap particle filter on the full dataset with the optimal ρ value. Present a plot of the estimated Infected, Exposed and Recovered states.

```
In [12]: optimal_rho = rho[mean == max(mean)]  
         model.set_param(rho = optimal_rho)
```

```
bpf3 = bpf(model = model, y = y_sthlm, numParticles = 200)

#R_true_bpf3 = population_size - np.sum(alpha[:3, trajectory, :], axis = 0)
R_bpf3 = population_size - np.sum(bpf3.alpha_filt[:3, 0, :], axis = 0)

fig, ax = plt.subplots(3, 1, figsize=(12, 8))

# Infected
ax[0].plot(bpf3.alpha_filt[2, 0, :], label='Filter Mean', color='orchid')
ax[0].set(title='Infected', xlabel='Time', ylabel='Number of Infected')
ax[0].legend()
ax[0].grid(True)

# Exposed
ax[1].plot(bpf3.alpha_filt[1, 0, :], label='Filter Mean', color='forestgreen')
ax[1].set(title='Exposed', xlabel='Time', ylabel='Number of Exposed')
ax[1].legend()
ax[1].grid(True)

# Recovered
ax[2].plot(R_bpf3, label='Filter Mean', color='royalblue')
ax[2].set(title='Recovered', xlabel='Time', ylabel='Number of Recovered')
ax[2].legend()
ax[2].grid(True)

plt.tight_layout()
plt.show()
```

